

Module 04

Your First Conversation



Meeting Your Agent for the First Time



Navigation Chart

By the end of this module, you will be able to:

1. **Navigate** the TUI (Terminal User Interface) confidently
2. **Introduce** yourself to your agent properly (the "brain dump")
3. **Name** your agent and give it an identity
4. **Understand** the token counter and why it matters
5. **Use** the essential slash commands: /compact , /status , /new , /help , /exit
6. **Explain** how context windows work and why conversations "fill up"
7. **Have** a productive first session with your agent



Ship's Logbook — Part 1

Term	Definition
TUI	Terminal User Interface — the interactive chat inside your terminal
Context window	The amount of text (in tokens) the AI can "remember" in a single conversation
Tokens	Units AI models use to process text. ~1 token = 3/4 of a word
Brain dump	Your initial download of information to your agent — who you are, your goals, your preferences



Ship's Logbook – Part 2

Term	Definition
Slash command	A command starting with / that instructs 🦀 OpenClaw to perform an action (not a message to the AI)
Session	A single conversation thread with your agent
Compact	The /compact command that summarizes your conversation to free up context window space

Opening the TUI



Launch the terminal chat interface:

```
openclaw tui
```

Key elements:

- **Top bar** — Model name and token count
- **Middle area** — The conversation (your messages + AI responses)
- **Bottom input** — Where you type messages and commands
- **Token counter** — Tokens used in this session (climbs with every message)

The Brain Dump – Why It Matters

The **brain dump** is the single most important thing you do in your first conversation.

- Without it, your agent is **generic** — it knows nothing about you
- With it, your agent becomes **yours** — personalized and useful
- Think of it as onboarding a new employee on their first day
- The more you share, the more useful your agent becomes
- It covers: who you are, your goals, your preferences, your rules

Hiring an employee and never telling them anything about the company or their role — that is what skipping the brain dump looks like.

Brain Dump Template (Part 1)

Copy, fill in, and paste as your first message:

Hi! I'd like to introduce myself and set you up as my personal AI assistant.

****Who I am:****

- Name: [Your name]
- Location: [City/Country]
- Occupation: [What you do]
- Background: [Brief career/life history]
- Expertise: [What you're good at]

****My personality and communication style:****

- I prefer [direct/detailed/casual/formal] communication
- I like [short answers / thorough explanations / both]
- I value [honesty/efficiency/creativity/precision]

Brain Dump Template (Part 2)

****My goals and ambitions:****

- Short-term (next 30 days): [Current focus]
- Medium-term (next 6 months): [What to accomplish]
- Long-term (1-3 years): [Bigger vision]

****What I want you to help me with:****

- [Scheduling, research, writing, coding, etc.]

****Rules for you:****

- [E.g., "Don't sugarcoat things"]
- [E.g., "Always ask before taking any action"]

Please remember all of this.

Alternative: Let Your Agent Interview You

I want you to interview me. Ask me questions to build my SOUL.md and USER.md files. One question at a time.

Often better -- the agent knows what format it needs.

Brain Dump Example

Who I am:

- Name: Alex
- Location: Austin, Texas
- Occupation: Freelance web developer
- Background: Self-taught, 3 years JS/React
- Expertise: Frontend dev, sales psychology

My goals:

- Short-term: Launch portfolio website this month
- Long-term: One-person business making \$10K/month

Rules:

- Don't sugarcoat things. Tell me the truth.
- Always ask before executing destructive commands.
- When I say "quick answer," give 1-2 sentences max.

Sending the Brain Dump

After you paste and send, your agent responds with a confirmation:

Thank you for the thorough introduction! I've noted all of this. A few highlights I want to confirm:

- I'll keep communication direct and honest
- I'll help you stay focused on priorities
- I'll check before any destructive actions

Is there anything else you'd like me to know?

- Review the response — does your agent understand your key points?
- Correct anything it misunderstood
- The brain dump is not permanent yet (that comes in Module 05)

Naming Your Agent

For this course, we use **your agent** — give it a name that resonates with you.

Send a message like:

From now on, your name is [Your Agent's Name]. You are my AI personal assistant. Embody these qualities: wisdom, patience, directness, and quiet confidence.

Choosing your own name:

- Pick something you enjoy saying — you will use it often
- Give it meaning (mythology, literature, history)
- Avoid generic names like "Assistant" or "Bot"
- The name becomes permanent in Module 05 when we edit `IDENTITY.md`

Essential Slash Commands



Command	What It Does	When to Use
/help	Lists all available commands	Need to find a command
/status	Shows model, tokens, gateway info	System overview
/compact	Summarizes conversation to free tokens	Token counter above 50K
/new	Starts a fresh session	Switching topics entirely
/think <level>	Sets thinking depth (off/low/medium/high)	Complex reasoning tasks
/exit	Leaves the TUI (also: Ctrl+C)	Done chatting

Understanding Tokens

What are tokens?

- The units AI models use to process text
- ~1 token = 3/4 of a word in English
- "Hello, how are you today?" = ~7 tokens
- A full page of text = ~500-800 tokens

Why tokens matter:

- **Cost** — You pay per token (input ~\$15/M, output ~\$75/M for Claude Opus 4.6)
- **Context window** — The model holds a finite number of tokens (~128K-200K); when full, old messages are dropped

Managing Token Usage

Strategy	What It Does	When to Use
/compact	Compresses conversation history	Token counter above 50,000
/new	Starts fresh session	Switching topics entirely
/think <level>	Sets AI thinking depth for complex tasks	Need deeper reasoning on hard problems
Short prompts	Fewer input tokens	When you do not need long explanations
"Quick answer" prefix	AI gives shorter responses	When you just need a fact
Cheaper model	Lower cost per token	Routine tasks, simple questions

Example: Token counter shows 80,000. You type /compact . It drops to ~27,000. Over 50,000 tokens of space freed.

Your First Real Conversations

Conversation 1: Test the Waters

Ask something you already know the answer to:

What are the three most important principles of personal finance?

Evaluate: Is it accurate? Does the tone match your brain dump preferences?

Conversation 2: Ask for a Plan

I want to learn Bitcoin in 30 days. Create a study plan –
30 minutes per day, complete beginner.

Tests: memory of your background, planning ability, personalization.

More First Conversations

Conversation 3: Test the Safety Boundary

Before you do anything, tell me: what files exist in my home directory? Don't list them – just tell me whether you CAN access them, and wait for my permission.

Tests: respect for "ask before acting," judgment about sensitive actions.

Conversation 4: Ask Your Agent to Reflect

What do you know about me so far? Summarize what you've learned and tell me the most important thing you can help me with.

Tests: how well your agent absorbed your brain dump.

The "Ask Before Acting" Habit

Add this to any request that involves action:

Before you do anything, give me a step-by-step plan of what you intend to do. Do not execute any commands until I approve the plan.

This creates a safety loop: agent explains the plan, you approve or modify, then agent acts.

Especially important for: file ops, system commands, network access, sending messages.

Voice Messages: Talk Instead of Type

Typing long messages is slow. **Telegram supports voice messages natively.**

Setup:

- Install **Groq** (free) for transcription: `openclaw skills install groq-transcription`
- Talk naturally into Telegram — your agent transcribes and processes

Best for:

- Brain dumps (talk about your life for 5 minutes instead of typing)
- Complex instructions (easier to explain by voice)
- Thinking out loud (let the agent capture your stream of consciousness)

Much faster for the initial brain dump. Just talk. Your agent listens.

► Shoals and Sandbars

Mistake	Fix
Skipping the brain dump	Take 15 minutes to fill in the template
Brain dump is too vague	Be specific: "help organizing my daily schedule" not "help with stuff"
Not using /compact	Compact when token counter passes 50,000
One endless conversation	Use /new for new topics
Treating the AI like Google	Use it for complex tasks: planning, writing, analyzing
Getting frustrated with wrong answers	LLMs hallucinate — verify important info
Not naming the agent	Name it now; you can always change it later

Hands on Deck: Write Your Brain Dump

Part 1: Write It (15 min)

- At least 5 facts about yourself
- At least 3 communication preferences
- At least 2 short-term and 2 long-term goals
- At least 3 specific use cases for help
- At least 2 rules for the agent

Part 2: Send It (5 min)

1. Open the TUI: `openclaw tui`
2. Paste your brain dump
3. Read your agent's response — does it show understanding?

Hands on Deck (continued)

Part 3: Test (10 min)

1. Ask your agent to **summarize** what it knows about you
2. Ask your agent to **suggest** one thing to focus on this week
3. Ask your agent to **explain its plan** before checking your Downloads folder

Part 4: Evaluate (5 min)

- Did your agent remember key details from your brain dump?
- Did the communication style match your preferences?
- Did your agent ask for permission before acting?
- What would you change about your brain dump?



Treasure Chest

1. **The brain dump is the foundation** — the more you share, the more useful your agent becomes
2. **Name your agent** — it makes the experience personal and reinforces the "employee" model
3. **Master five commands:** /help , /status , /compact , /new , /exit
4. **Watch the token counter** — compact at 50K, new session when switching topics
5. **Always ask for a plan before action** — your safety net
6. **Test gradually** — questions and planning before executing commands
7. **Brain dump is not permanent yet** — Module 05 encodes it into persistent 🧫 memory files
8. **LLMs hallucinate** — verify important information

Next Port of Call

Module 05: Workspace and Memory

Your brain dump lives in one session. In Module 05, we turn it into permanent configuration that persists across every session, every restart, and every update.

This is where your agent goes from "chatbot" to "AI employee who truly knows you."