

# Module 09



## Advanced Concepts

Cost Optimization, Model Routing, and Power-User Techniques



# Navigation Chart

By the end of this module, you will be able to:

1. **Understand** why 🦀 OpenClaw burns money by default and how to fix it
2. **Implement** the "Brains and Muscles" model for cost optimization
3. **Set up** intelligent model routing (manual, ClawRouter, or OpenRouter)
4. **Enable** prompt caching for up to 90% savings on system prompts
5. **Switch** between AI models for different tasks
6. **Understand** 💭 multi-agent setups
7. **Use** reverse prompting to let the AI guide your decisions



# Ship's Logbook (Part 1)

| Term               | Definition   |
|--------------------|--|
| Brains and Muscles | Use an expensive model (the "brain") for thinking, and cheaper models (the "muscles") for execution                |
| Model switching    | Changing which AI model handles a conversation mid-stream  |
| ClawRouter         | LOBster 🦀 OpenClaw-native routing tool that classifies request complexity and routes to the cheapest capable model |
| Prompt caching     | Provider feature that remembers static parts of your prompt between calls so you pay less                          |



## Ship's Logbook (Part 2)

| Term                 | Definition  |
|----------------------|---|
| Context accumulation | Session history growing with every message -- a mature session can reach 200K+ tokens |
| Multi-agent          | Running more than one 🦐 OpenClaw agent for different life domains                     |
| Reverse prompting    | Instead of telling the AI what to do, you ask the AI what YOU should do               |
| Open Router          | A service providing access to 300+ AI models through a single API                     |
| Agentic company      | Structuring your AI agents like employees with different roles                        |

# The Problem: Why 🦀 OpenClaw Burns Money

Running 🦀 OpenClaw with a single frontier model for all tasks is the **#1 cost mistake**.

| Cost Mechanism             | What Happens   |
|----------------------------|--|
| Context accumulation       | Session history grows to 200K+ tokens; every follow-up carries enormous overhead |
| System prompt re-injection | 3,000-14,000 tokens re-sent with every API call                                  |
| Tool output storage        | File listings, browser snapshots bloat context over time                         |
| Heartbeat overhead         | 48 full-context calls/day on Opus = \$70-200/month                               |
| Cron job overhead          | Each trigger creates a fresh conversation with full context                      |

*90% of what your agent does is routine work that doesn't need a \$5/MTok model.*

# The Brains and Muscles Model

---

Split your workload between expensive and cheap models:

## The Brain (Opus 4.6)

- Complex decisions and strategy
- Creative writing
- Prompt injection defense
- Orchestrating other models

## The Muscles (Cheaper Models)

- **Haiku 4.5** -- quick lookups, heartbeats (\$1-5/MTok)
- **Sonnet 4.5** -- code generation, summarization (\$5-15/MTok)
- **Codex** -- development tasks (varies)

# The Opus Orchestra Pattern

A named strategy used by experienced 🦐 OpenClaw operators:

| Role                    | Model                     | Cost     |
|-------------------------|---------------------------|----------|
| Orchestration/decisions | Opus 4.6                  | Premium  |
| Sub-agent execution     | Sonnet 4.5                | Mid-tier |
| Heartbeats/routine      | Haiku 4.5 or Gemini Flash | Minimal  |

## Real Cost Benchmarks (Community Data)

- **Without optimization:** \$100-400/month
- **Well-configured:** \$10-50/month
- **Extreme optimization:** <\$10/month
- **One user's \$800 mistake** -- 8 agents, no limits, pay-per-use

# Practical Model Assignments

| Task                    | Recommended Model   | Why                          |
|-------------------------|---------------------|------------------------------|
| Important conversations | Opus 4.6            | Best reasoning, most nuanced |
| Code generation         | Sonnet 4.5 or Codex | Optimized for code, cheaper  |
| Heartbeats              | Haiku 4.5           | Routine checks, very cheap   |
| Web search              | Perplexity Pro      | Built for search             |
| Document summarization  | Sonnet 4.5          | Good enough, much cheaper    |
| Quick factual lookups   | Haiku 4.5           | Fast and cheap               |

# Practical Model Assignments (continued)

| Task                     | Recommended Model                 | Why  |
|--------------------------|-----------------------------------|--|
| Creative writing         | Opus 4.6                          | Best quality   |
| Security-sensitive tasks | Opus 4.6                          | Best prompt injection resistance                               |
| Budget tasks             | Kimi K2.5 (free via NVIDIA)       | Comparable to Opus, zero cost                                  |
| Ultra-cheap tasks        | MiniMax M2.5 (~\$10-50/mo)        | Surprisingly capable; Max plan gives 1,000 prompts per 5 hours |
| Free tier exploration    | Google Cloud (\$300 free credits) | Gemini models with sign-up credits (rate limits apply)         |

# Model Switching in the TUI

Switch models during a conversation without starting a new session:

```
/model claude-opus-4-6      # Expensive, best quality  
/model claude-sonnet-4-5     # Good balance  
/model claude-haiku-4-5      # Fast and cheap
```

## When to switch:

- **Opus:** Important decisions, untrusted input, nuanced conversations
- **Sonnet:** Code generation, summarization, drafts
- **Haiku:** Quick facts, formatting, template generation, testing

# Intelligent Routing: ClawRouter

ClawRouter is the hottest tool in the 🦀 OpenClaw ecosystem -- 2,400 GitHub stars in its first 11 days.

## How it works:

Analyzes each query using a lightweight classifier, then routes to the cheapest capable model.

| Tier    | Complexity                          | Routed To              | Cost              |
|---------|-------------------------------------|------------------------|-------------------|
| Simple  | Basic lookups, acknowledgments      | DeepSeek, Gemini Flash | ~\$0.27-0.60/MTok |
| Medium  | Moderate tasks, summarization       | GPT-4o-mini, Sonnet    | ~\$1-5/MTok       |
| Complex | Analysis, writing, decisions        | Claude Sonnet          | ~\$3-15/MTok      |
| Heavy   | Multi-step reasoning, agentic tasks | Opus 4.6, Kimi K2.5    | ~\$5-25/MTok      |

# ClawRouter Profiles

| Profile | Description  |
|---------|--|
| Auto    | Balanced quality and cost (recommended starting point) |
| Eco     | Maximum savings, up to 95-100% on simple queries       |
| Premium | Best quality, less aggressive routing                  |
| Free    | Zero-cost models only (limited capability)             |

## Other routing options:

| Approach                  | Best For   |
|---------------------------|--|
| Manual rules in AGENTS.md | Getting started, simple setups                       |
| ClawRouter                | Most users -- best balance of automation and control |
| Custom routing 🐟 skill    | Power users who want exact control                   |
| OpenRouter auto-routing   | Users who want zero configuration                    |

# Prompt Caching: The Hidden Savings Multiplier

Every API call re-sends your full system prompt (SOUL.md, AGENTS.md,  MEMORY.md) -- typically **3,000-14,000 tokens**. You pay full price every time.

## With caching enabled:

- First call pays full price
- Subsequent calls within the cache window pay **90% less** for cached tokens

## Enable it:

```
{  
  "cacheRetention": "long",  
  "cacheSystemPrompts": true,  
  "cacheThresholdTokens": 2048  
}
```

- "Long" uses Anthropic's extended cache (~55 minutes)

# The 55-Minute Heartbeat Trick (Revisited)

Combine three optimizations for maximum savings:

1. Set heartbeat interval to **55 minutes** (stays within cache window)
2. Enable **prompt caching**
3. Route heartbeats to **Haiku 4.5**

| Without Optimization          | With Optimization              |
|-------------------------------|--------------------------------|
| Opus for heartbeats           | Haiku for heartbeats           |
| Full system prompt every call | Cached system prompt (90% off) |
| <b>~\$100+/month</b>          | <b>~\$0.50/month</b>           |

**99.5% reduction** on heartbeat costs.

*Prompt caching also works for regular conversations -- 90% off system prompt costs for every*

# Upgrading Web Search with Perplexity Pro

LOBSTER OpenClaw's built-in search is functional, but **Perplexity Pro** is significantly better for research.

## Setup via Open Router:

1. Create an account at [openrouter.ai](https://openrouter.ai)
2. Get your 🔑 API key
3. Add to 🦀 OpenClaw:

```
openclaw config provider add openrouter --key YOUR_KEY  
openclaw config search model perplexity-pro
```

## Cost tip:

Save research results as markdown files so you do not pay for the same search twice:

# 🐙 Agent-to-Agent Communication (`sessions_*` Tools)

Three built-in tools for agent communication -- no external service required:

| Tool                          | What It Does                                 |
|-------------------------------|--|
| <code>sessions_list</code>    | Discover active sessions and their metadata  |
| <code>sessions_history</code> | Fetch transcript logs from another session   |
| <code>sessions_send</code>    | Message another session; supports reply-back |

**Example flow:** Chief of Staff receives email about crypto → uses `sessions_send` to forward to Crypto Trader → Trader evaluates and sends back assessment → Chief of Staff summarizes for you on Telegram.

Individual agents communicating through `sessions_send` is more reliable and cheaper than one mega-agent doing everything.



# Multi-Agent Management

As you get comfortable, you might want separate agents for different areas of your life.

## The Agentic Company Structure:

Real example (7 days, ~\$600):

- **Sam (Chief of Staff)** -- email, calendar, CRM
- **Midas (Crypto Trader)** -- autonomous portfolio management
- **Ritam (Physics Research)** -- cross-domain synthesis
- **Personal assistant** -- daily life management

## Individual Agents > Sub-Agents

Community consensus: **multiple individual agents** with separate Telegram bots work better than one agent spinning up sub-agents. Sub-agents lose context and forget. Individual agents maintain their own  memory.

# ► Rough Waters: The Coordination Tax

---

Google DeepMind research: accuracy **saturates or degrades past 4 agents** ("Coordination Tax")

The "17x error trap": more agents = multiplied error rate, not throughput.

## The 17 → 4 Consolidation

One user went from 17 agents to 4 core roles:

1. **Architect** (CEO) -- strategy, priorities
2. **Builder** (CTO) -- engineering, quality
3. **Money Maker** -- growth, pricing, channels
4. **Operator** (COO) -- processes, tools, financial ops

Everything else: a **specialist library** (36+ types) spawned on demand by core agents.

# The Freshman Rule & 🐙 Multi-Agent 🍄 Memory

## The Freshman Rule

- One task at a time -- do not stack requests
- Ground-up instructions every time -- assume the agent knows nothing unless it is in files
- If it is not in files, it does not exist -- never assume agents remember past sessions

## Multi-Agent Memory: 4-Layer Architecture

| Layer          | What  | Example                                |
|----------------|---|--|
| 1: Private     | Each agent has its own <code>MEMORY.md</code>         | Midas knows crypto; Sam knows email    |
| 2: Shared dir  | A common <code>_shared/</code> folder all agents read | User profile, roster, conventions      |
| 3: QMD paths   | Cross-agent search via shared indexed paths           | Any agent searches another's docs      |
| 4: Coordinator | Dedicated agent maintaining consistency               | Resolves conflicts, propagates updates |

# Reverse Prompting

## The normal way (You tell AI what to do):

Build me a landing page for my freelance business.

## The reverse way (AI tells YOU what to do):

Based on everything you know about me, my goals, and my current situation, what should I be working on right now? What would move the needle the most?

## Why it is powerful:

- your agent has **perfect 🌸 memory** of every goal and conversation
- your agent has **no ego** -- will tell you uncomfortable truths
- your agent **sees patterns** -- spots contradictions in your behavior

# Reverse Prompting: Try These

## Strategic:

What am I overlooking? Based on my goals and what I've been working on, what blind spots do you see?

## Accountability:

Review my goals from last month. What did I accomplish? What did I skip? What pattern do you notice?

## Self-improvement:

Based on our interactions, what are my three biggest strengths and three biggest weaknesses in how I work?

# Self-Improvement Workflows

---

## Weekly Agent Tuning (make it a cron job):

Ask your agent:

1. What instructions in your core files feel unclear?
2. What information about me seems outdated?
3. What do I keep asking that should be in your 🌸 memory?
4. What 🐟 skills or automations would make us more efficient?
5. What suggestions do you have for improving our workflow?

Then **actually implement** the suggestions by editing core files.

## Context Optimization:

- Use `/compact` to reduce conversation history
- Start new sessions with `/new` for new topics

# ► Shoals and Sandbars: Chat Quality ≠ Agent Quality

A model that writes beautiful essays may **completely fail** at agentic work. Chat benchmarks do not measure tool-calling reliability, multi-step planning, or error recovery.

## Proven for Agentic Work

| Model                        | Strengths   |
|------------------------------|---|
| Claude Sonnet 4.5 / Opus 4.6 | Reliable tool calls, strong planning, consistent execution        |
| GPT-5.2                      | Solid multi-step reasoning, dependable function calling           |
| Kimi K2 / K2.5               | Comparable to Opus on many tasks, excellent cost-to-quality ratio |

## Avoid for Autonomous Tasks

| Model             | Problem  |
|-------------------|--|
| DeepSeek Reasoner | Great at thinking and analysis, but <b>broken tool calls</b> -- it reasons beautifully then fails to execute |

# 🚩 Shoals and Sandbars: Common Mistakes

| Mistake                                      | Fix   |
|--|---|
| Using Opus for everything                    | Set up brains and muscles model                               |
| Not enabling prompt caching                  | Enable <code>cacheSystemPrompts</code> -- instant 90% savings |
| Heartbeat interval at 30 min without caching | Set to 55 minutes to stay in cache window                     |
| Setting up 5 agents before mastering 1       | Master one agent, then expand                                 |
| Ignoring reverse prompting                   | Try it once a week  |
| Never doing agent tuning                     | Schedule weekly tuning (10 min)                               |
| Expecting local models to match Opus         | Use local for low-stakes tasks, cloud for important work      |

# ⚙️ Hands on Deck: Design Your Ideal 🧑 Multi-Agent Setup

---

## Part 1: Map Your Life Domains (10 min)

- List 3-5 areas: personal, professional, business, creative, other

## Part 2: Design the Roles (10 min)

For each domain:

- What would the agent focus on?
- What personality should it have?
- What model should it use?

## Part 3: Start Simple (5 min)

- Pick the **ONE** domain that would benefit most
- That is your first agent – master it before adding more



# Treasure Chest

1. 🦀 **OpenClaw burns money by default** -- five cost mechanisms compound to create huge bills
2. **Brains and Muscles saves serious money** -- expensive model for thinking, cheap for execution
3. **Prompt caching is the biggest hidden savings** -- 90% off system prompt costs
4. **ClawRouter automates routing** -- classifies complexity, routes to cheapest capable model
5. **Model switching is easy** -- `/model [name]` changes mid-conversation
6. **Perplexity Pro is worth the upgrade** -- significantly better research
7. 💫 **Multi-agent is powerful but complex** -- master one agent first
8. **Reverse prompting is the secret weapon** -- let your agent analyze your situation

# Next Port of Call

## Module 10: Security Hardening

*Sandboxing with Docker, tool policies,  gateway hardening, and incident response. Time to lock everything down.*